

- Elasticsearch Data Modeling
- Elasticsearch Data Processing
- Elasticsearch from Dev to Production
- Elasticsearch Cluster Deployment
- Elasticsearch Nodes and Index Management
- Elasticsearch Advanced Tips and Tricks

Module 3

# Elasticsearch from Dev to Production



# Topics

- Securing Elasticsearch
- Development vs. Production Mode
- Scaling Elasticsearch
- Server Configuration Best Practices



Elasticsearch From Dev to Production

Lesson 1

# Securing Elasticsearch




# Security Considerations

- Make sure to secure your cluster before going live:
  - firewalls
  - reverse proxy
  - *Elastic Security*


# Firewall

- Add ***firewall rules*** with IP filtering to only allow access from the machines running your application
  - never expose port 9200 http publicly on the internet
  - restrict access of ports to only systems necessary
  - be careful, don't block port 9300 between the nodes



# Reverse Proxy

- You can put Elasticsearch behind a *reverse proxy*
  - gain basic authentication and authorization with OSS version
  - Elastic security (Basic license) is recommended instead



# Elastic Security

- The Elasticsearch Basic subscription (free) includes **core security features** such as:
  - encrypted communications
  - role-based access control
  - and more...
- It enables the community to encrypt network traffic, create users and define roles to protect their data
  - <https://www.elastic.co/blog/security-for-elasticsearch-is-now-free>

# Configuring Elastic Security

- Enable Elastic security feature on every node (`elasticsearch.yml`)

```
xpack.security.enabled: true
```

- Configure Transport Layer Security (TLS/SSL) for internode-communication
- If you plan to run Elasticsearch in a Federal Information Processing Standard (FIPS) 140-2 enabled JVM

```
xpack.security.fips_mode.enabled: true
```

- Restart Elasticsearch on each node of your cluster
- Set the passwords for all built-in users

```
./elasticsearch/bin/elasticsearch-setup-passwords interactive
```

# Configuring Transport Layer Security

- Generate a certificate and private key for each node in your cluster

```
./elasticsearch/bin/elasticsearch-certutil cert -out config/elastic-certificates.p12
```

- Copy the PKCS#12 keystore to each node in your cluster
- Enable TLS communication on each node in your cluster (**elasticsearch.yml**)

```
xpack.security.transport.ssl.enabled: true  
xpack.security.transport.ssl.verification_mode: certificate  
xpack.security.transport.ssl.keystore.path: elastic-certificates.p12  
xpack.security.transport.ssl.truststore.path: elastic-certificates.p12
```

# Configuring Kibana After Enabling Security

- Configure Kibana to use the appropriate built-in user (**kibana.yml**)

```
elasticsearch.username: "kibana"  
elasticsearch.password: "kibanapassword"
```

- You can also use Kibana keystore
  - create a keystore to avoid having passwords in clear text in the configuration file

```
bin/kibana-keystore create
```


- add sensitive information to your keystore

```
bin/kibana-keystore add elasticsearch.username  
kibana  
bin/kibana-keystore add elasticsearch.password  
kibanapassword
```

- those parameters don't need to be defined in **kibana.yml** anymore


# After Configuring Elastic Security

- Login with user and password that you defined before



# Defining Users and Roles

- You can use Kibana to set up roles and users to control access to Elasticsearch



The screenshot shows the 'Roles' page in Kibana. On the left, there's a sidebar with a shield icon labeled 'Security', followed by 'Users' and 'Roles' (which is underlined, indicating it's the active tab). The main area is titled 'Roles' with the subtitle 'Apply roles to groups of users and manage permissions across the stack'. It features a search bar and a 'Create role' button. A table lists 33 predefined roles, each with a checkmark in the 'Reserved' column:

Role ↑	Reserved ⓘ
apm_system	✓
apm_user	✓
beats_admin	✓
beats_system	✓
blogs_READONLY	
code_admin	✓
code_user	✓
ingest_admin	✓
kibana_dashboard_only_user	✓
kibana_system	✓
kibana_user	✓
logstash_admin	✓
logstash_system	✓
machine_learning_admin	✓
machine_learning_user	✓
monitoring_user	✓
remote_monitoring_agent	✓
remote_monitoring_collector	✓
reporting_user	✓
rollup_admin	✓
rollup_user	✓
snapshot_user	✓
superuser	✓
transport_client	✓
watcher_admin	✓
watcher_user	✓

# Defining Roles and Privileges

- Roles and privileges can be defined at a very fine-grained level

The screenshot shows the Elasticsearch Role Management interface. It includes sections for Cluster privileges, Run As privileges, and Index privileges.

**Cluster privileges:** Manage the actions this role can perform against your cluster. Options include: monitor, manage, manage\_index\_templates, manage\_ml, monitor\_ml, manage\_watcher, and monitor\_watcher.

**Run As privileges:** Allow requests to be submitted on the behalf of other users. A dropdown menu is shown with the placeholder "Add a user...".

**Index privileges:** Control access to the data in your cluster. Options include: my\_index\*, manage, monitor, read, and read\_cross\_cluster.

**Add index privilege:** A button to add more index privileges.

# Defining Roles and Privileges to Users

- Add users with role-based access control

Security

[Users](#)

Roles

Users

<input type="checkbox"/> Full Name ↑	User Name	Email Address	Roles	Reserved
<input type="checkbox"/> Blog User	blogs_user	nobody@elastic.co	blogs_READONLY, kibana_user	
<input type="checkbox"/>	elastic		superuser	✓
<input type="checkbox"/>	kibana		kibana_system	✓
<input type="checkbox"/>	logstash_system		logstash_system	✓
<input type="checkbox"/>	beats_system		beats_system	✓
<input type="checkbox"/>	apm_system		apm_system	✓
<input type="checkbox"/>	remote_monitoring_user		remote_monitoring_collector, remote_monitoring_agent	✓

Rows per page: 20 ▾

Create new user



Elasticsearch From Dev to Production

Lesson 1

# Review - Securing Elasticsearch



# Summary

- You can secure your cluster with firewalls, reverse proxy, and Elastic Security
- ***Elastic Security*** provides a complete solution for securing Elasticsearch
- If **Elastic Security** is enabled, unless you have a trial license, you must configure SSL/TLS for internode-communication
- **Make sure to secure your cluster!**

# Quiz


1. Explain three options to secure your cluster.
2. **True or False.** An Elasticsearch cluster comes with security enabled by default.
3. **True or False.** It is required to enable SSL/TLS communication when Elastic security is enabled on the basic license.



Elasticsearch From Dev to Production

Lesson 1

# Lab - Securing Elasticsearch





Elasticsearch From Dev to Production

Lesson 2

# Development vs. Production Mode




# When You Go To Production

- You might want to care about
  - development vs. production mode
  - caching
  - Linux packages

# Development vs. Production Mode


# HTTP vs. Transport

- There are two important network communication mechanisms in Elasticsearch to understand
  - **HTTP**: address and port to bind to for HTTP communication, which is how the Elasticsearch REST APIs are exposed
  - **transport**: used for internal communication between nodes within the cluster
- The defaults are fine for downloading and playing with Elasticsearch, but not useful for production systems
  - bind to **localhost** by default



# Development vs. Production Mode

- Every Elasticsearch instance is either in development mode or production mode
  - **development mode**: if it does *not* bind transport to an external interface (the default)
  - **production mode**: if it does bind transport to an external interface



# Bootstrap Checks

- Elasticsearch has ***bootstrap checks*** upon startup
  - they inspect a variety of Elasticsearch and system settings,
  - and compare them to values that are safe for the operation of Elasticsearch
- Bootstrap checks behave differently depending on the mode
  - **development mode**: any bootstrap checks that fail appear as ***warnings*** in the Elasticsearch log
  - **production mode**: any bootstrap checks that fail will cause Elasticsearch to ***refuse to start***
  - <https://www.elastic.co/guide/en/elasticsearch/reference/current/bootstrap-checks.html>

# Bootstrap Checks

- A node in *production mode* must pass all of the checks, or the node will not start
  - the bootstrap checks fit into two categories

JVM Checks	Linux Checks
<input checked="" type="checkbox"/> heap size	<input checked="" type="checkbox"/> file descriptor
<input checked="" type="checkbox"/> memory lock	<input checked="" type="checkbox"/> maximum number of threads
<input checked="" type="checkbox"/> server JVM	<input checked="" type="checkbox"/> max file size
<input checked="" type="checkbox"/> not use serial collector	<input checked="" type="checkbox"/> maximum size virtual memory
<input checked="" type="checkbox"/> OnError and OnOutOfMemoryError	<input checked="" type="checkbox"/> maximum map count
<input checked="" type="checkbox"/> not use early-access snapshots	<input checked="" type="checkbox"/> system call filter
<input checked="" type="checkbox"/> not use G1GC collector	<input checked="" type="checkbox"/> java security permissions
<input checked="" type="checkbox"/> disable swapping	<input checked="" type="checkbox"/> discovery configuration



# Caching

# Node Query Cache

- **Caches search results** for non-scoring queries
  - uses Least Recently Used (LRU) eviction policy
  - per node, shared by all the shards on that node
- **Use filters whenever possible:**

```
GET blogs/_search
{
  "query": {
    "range": {
      "publish_date": {
        "gte": "2016-01-01",
        "lte": "2018-12-31"
      }
    }
  }
}
```

Cannot be cached

```
GET blogs/_search
{
  "query": {
    "bool": {
      "filter": {
        "range": {
          "publish_date": {
            "gte": 2016-01-01,
            "lte": 2018-12-31
          }
        }
      }
    }
  }
}
```

Can be cached

# Shard Request Cache

- Caches shard-level *aggregation results, counts and suggestions* by default
  - uses Least Recently Used (LRU) eviction policy
  - per node, shared by all the shards on that node
  - uses the JSON request body as the cache key
- Only if the data has not changed and **size** is set to 0

```
GET logs_server*/_search
{
  "aggs": {
    "logs_by_day": {
      "date_histogram": {
        "field": "@timestamp",
        "interval": "day"
      }
    }
  }
}
```

Is not cached

```
GET logs_server*/_search
{
  "size": 0,
  "aggs": {
    "logs_by_day": {
      "date_histogram": {
        "field": "@timestamp",
        "interval": "day"
      }
    }
  }
}
```

Is cached

# Shard Request Cache

- Can also be used to cache *hits*
- Use the **request\_cache** parameter to force caching:

```
GET blogs/_search?request_cache=true
{
  "query": {
    "match": {
      "content": "filebeat"
    }
  },
  "aggs": {
    "top_categories": {
      "terms": {
        "field": "category.keyword",
        "size": 10
      }
    }
  }
}
```

caches even when size > 0

# Linux Packages

# Linux Packages

- The tarball package is good for testing and development
- On a production environment you should take advantage of your package management system
  - automatic dependency verification
  - better directory infrastructure
  - binaries and libraries on standard locations
  - configuration and log files on standard locations
  - easier to update through package management systems
  - easier to start and stop through daemons
  - automatic startup on machine reboot

# Configuration Management

- It becomes difficult to manage your Elasticsearch settings as you grow your cluster
  - it is really hard to edit many configuration files by hand without making a mistake
- It is a good idea to use a configuration management tool
  - for example: Puppet, Chef, Ansible
- Configuration management tools help make your cluster consistent by automating the process of config changes
  - when you install Elasticsearch through Linux packages it is a bit easier to automate your configuration process
  - all files are in standard locations and an update should not break your automation process





Elasticsearch From Dev to Production

Lesson 2

# Review - Development vs. Production Mode



# Summary

- A node in ***production mode*** must pass a series of checks, or the node will not start
- A query in a filter context can be ***cached*** by Elasticsearch to improve performance
- Installing Elasticsearch with Linux packages makes it easier to manage your cluster configurations

# Quiz


- 1. True or False.** In production mode, if a bootstrap check fails then the node will not start.
- 2. True or False.** Elasticsearch caches any query to improve performance.
3. Why would you use Linux packages to install Elasticsearch?



Elasticsearch From Dev to Production

Lesson 2

# Lab - Development vs. Production Mode





Elasticsearch Advanced Tips and Tricks


Lesson 3

# Scaling Elasticsearch




# Designing for Scale

- Elasticsearch is built to scale
  - and the default settings can take you a long way
- Proper design can make scaling easier



# One Shard...

- ...does not scale very well:




Adding a second node would not provide any scaling for **my\_index**

```
PUT my_index
{
  "settings": {
    "number_of_shards": 1,
    "number_of_replicas": 0
  }
}
```

# Two Shards...


- ...can scale if we add a node:



```
PUT my_index
{
  "settings": {
    "number_of_shards": 2,
    "number_of_replicas": 0
  }
}
```

# Balancing of Shards


- Elasticsearch automatically balances the shards:
  - node1 is now responsible for half the amount of data
  - write throughput has doubled (twice the disk IO available)
  - the memory pressure on node1 is less than before
  - when searches are executed, we now use the resources of both node1 and node2



# Shard Overallocation


- If you are expecting your cluster to grow, then it is good to plan for that by overallocating shards:
  - $\#shards > \#nodes$
- Shards can move between nodes quickly as the cluster grows
  - and there is no downtime during shard relocation
- Note: this strategy works for static, but not time-series data
  - for time-series data, you will have multiple indices with few shards

```
PUT my_index
{
  "settings": {
    "number_of_shards": 6,
    "number_of_replicas": 0
  }
}
```



# “Too Much” Overallocation

- A little overallocation is good
- A “kagillion” shards is not:
  - each shard comes at a cost (Lucene indices, file descriptors, memory, CPU)
  - also, a search request needs to hit every shard in the index
- A shard typically can hold at least 10s of gigabytes
  - depends on the use case
  - a 1GB shard is sub-utilized



# Do not Overshard

- Business Requirements

- 1GB per day
  - 6 months retention
  - **~180GB**




we could easily have this data in 10 shards

- Common Scenario

- 3 different logs
  - 1 index per day each
  - 5 shards (default before 7.0)
  - 6 months retention
  - **~2700 shards**

too many shards for no good reason!


## Cluster *my\_cluster*



# Scaling for Reads

# Scaling for Reads

- Searches and aggregations scale with replicas
- For example, one primary and as many replicas as you have additional nodes
  - use **auto\_expand\_replicas** setting to change the number of replicas automatically, as you add/remove nodes



# Optimizing for Read Throughput


- Avoid **nested types** and the **join** datatype
  - create flat, denormalized documents
- Query the smallest number of fields
  - consider **copy\_to** over **multi\_match**
- Map identifiers as **keyword** instead of as a number
  - term queries on keyword fields are very fast
- Force-merge read-only indices
- Limit the scope of aggregations (upcoming topic!)
- Use filters, as they are cacheable



# Scaling for Writes

# Scaling for Writes

- Write throughput scales by increasing number of primaries
- Configuring more primary shards allows Elasticsearch to "fan out" the writes, so each shard is doing less work
- Maximize throughput by using disks on all machines
- When an index is done with writes, you can shrink it



# Optimizing for Write Throughput

- Use `_bulk` API to minimize the overhead of HTTP requests
- Parallelize your write requests
- Disable refreshing every second
  - set `index.refresh_interval` to `-1` for very large writes (then back to default when finished indexing)
  - set `index.refresh_interval` to `30s` to increase indexing speed but affect search as little as possible (e.g. logs use case)
- Disable replicas, then re-enable after very large writes
  - every document also needs to be written to every replica
- Use auto-generated IDs
  - Elasticsearch won't check whether a doc ID already exists





Elasticsearch Advanced Tips and Tricks

Lesson 3

# Review - Scaling Elasticsearch



# Summary

- If you are expecting your cluster to grow, then it is good to plan for that by ***overallocating*** shards
- A little overallocation is good. A “kagillion” shards is not
- You can scale the read workload of your cluster by adding more nodes and ***increasing the number of replicas*** of your indices
- You can scale the write workload of your cluster by adding more nodes and ***increasing the number of primaries*** of your indices

# Quiz


1. If you have a two node cluster, why would you ever create an index with more than two primary shards?
2. **True or False.** To maximize read throughput you need to divide your data over as many primaries as possible
3. **True or False.** Using auto-generated IDs is more efficient than providing your own IDs



Elasticsearch Advanced Tips and Tricks

Lesson 3

# Lab - Scaling Elasticsearch





Elasticsearch Advanced Tips and Tricks

Lesson 4

# Server Configuration Best Practices



# Server Configuration Best Practices

- Hardware Best Practices
- JVM Settings

# Hardware Best Practices

# Networking Best Practices

- Avoid running over WAN links between data centers
  - not officially supported by Elastic
- Try to have zero (or very few) hops between nodes
- If you have multiple network cards
  - separate **transport** and **http** traffic (transport on faster NIC)
  - bind to different network interfaces
  - use separate firewall rules for each kind of traffic
- Use long-lived HTTP connections
  - client libraries support this
  - or use a proxy/load-balancer

# Storage Best Practices

- Prefer solid state disks (SSDs)
  - segments are immutable
  - so the *write amplification factor* approaches one and is a non-issue
- Local disk is king!
  - in other words, avoid NFS, SMB, AWS EFS, or Azure filesystem
- Elasticsearch does not need redundant storage
  - replicas = software provided HA
  - local disks are better than **SAN**
  - **RAID1/5/10** is not necessary



# Storage Best Practices

- If you have multiple disks in the same server
  - you can set **RAID 0** or **path.data**
- **RAID 0**
  - splits ("stripes") data evenly across two or more disks
  - perfect distribution of the data across the disks
  - However, if you lose one disk, you lose the data on all disks
- **path.data**
  - allows you to distribute your index across multiple disks
  - potential to have an unbalanced distribution (all files belonging to a shard will be stored on the same data path)
  - if you lose one disk, the shards on all other disks are preserved
    - may generate node level watermark issues if disks have different sizes



# Storage Best Practices

- SSDs: Use **noop** or deadline scheduler in the OS
  - for more detail, see:  
[https://www.elastic.co/guide/en/elasticsearch/guide/current/hardware.html#\\_disks](https://www.elastic.co/guide/en/elasticsearch/guide/current/hardware.html#_disks)

```
echo noop > /sys/block/{DEVICE}/queue/scheduler
```
- SSDs: Trim your disks periodically:
  - <https://www.elastic.co/blog/is-your-elasticsearch-trimmed>
- Spinning disks are OK for **warm** nodes
  - but don't forget to disable concurrent merges

```
index.merge.scheduler.max_thread_count: 1
```

# Hardware Selection

- Elasticsearch was designed to run on commodity hardware
- In general, choose ***large*** machines over ***x-large*** machines
  - loss of a large node has a greater impact
  - ***prefer*** 6 machines with 4-8cpu x 64gb x 4 1tb drives
  - ***avoid*** 2 machines with 12cpu x 256gb x 12 1tb drives
- Avoid running multiple nodes on one server
  - one Elasticsearch instance can fully consume a machine
  - resource contention + multiple JVMs can be problematic
- Larger machines can be helpful as ***warm*** or ***cold*** nodes
  - configure shard allocation filtering

# Hardware Selection

- 4 CPU
- 64 GB RAM
- 4 x 1TB disks

- 4 CPU
- 64 GB RAM
- 4 x 1TB disks


- 4 CPU
- 64 GB RAM
- 4 x 1TB disks

- 4 CPU
- 64 GB RAM
- 4 x 1TB disks



- 12 CPU
- 256 GB RAM
- 12 x 1TB disks

- 12 CPU
- 256 GB RAM
- 12 x 1TB disks



# Cloud Strategies

- On Cloud, use the *discovery plugin*
  - because IPs can change frequently, the plugin dynamically configures the discovery seeds
- Span the cluster across more than one AZ
  - use shard awareness and forced awareness
- Prefer ephemeral storage over network storage
  - as noted earlier, local SSD-backed storage is preferred
- Snapshot to cloud storage with the **repository plugins**
- Avoid instances marked with low networking performance
  - network latency will kill performance



# Throttles

- Elasticsearch has relocation and recovery throttles
  - it ensures that these tasks do not have a negative impact
  - you can speed up the process by throttling up
- Recovery
  - for faster recovery, temporarily increase the number of concurrent recoveries

```
PUT _cluster/settings
{
  "transient": {
    "cluster.routing.allocation.node_concurrent_recoveries": 2
  }
}
```

- Relocation
  - for faster rebalancing of shards, increase

```
"cluster.routing.allocation.cluster_concurrent_rebalance" : 2
```

# JVM Settings

# JVM Configuration

- Since Elasticsearch 6.0, only 64-bit JVMs are supported
- You can configure the Java Virtual Machine (JVM) in two ways:
  - the **config/jvm.options** file (preferred)

```
-Xms30g  
-Xmx30g
```



- setting the **ES\_JAVA\_OPTS** environment variable

```
ES_JAVA_OPTS="-Xms30g -Xmx30g" bin/elasticsearch
```

- Elasticsearch has very good JVM defaults
  - avoid and be careful when changing them

# What Goes on the Heap?

- Some of the major usage of the heap by Elasticsearch includes:



# JVM Heap Size

- By default, the JVM heap size is 1 GB
  - likely not high enough for production
  - you can change it using **Xms** (min heap) and **Xmx** (max heap)

```
ES_JAVA_OPTS="-Xms8g -Xmx8g" ./bin/elasticsearch
```

- Some guidelines for configuring the heap size
  - set **Xms** and **Xmx** to the same size (bootstrap check)
  - set **Xmx** to no more than 50% of your physical RAM
  - do not exceed more than 30GB of memory (to not exceed the compressed ordinary object pointers limit)
- Leave as much memory to the filesystem cache as possible
- <https://www.elastic.co/blog/a-heap-of-trouble>

# Production JVM Settings

- JDKs have two modes of a JVM: client and server
  - server JVM is required in production mode
- Configure the JVM to disable swapping
  - by requesting the JVM to lock the heap in memory through mlockall (Unix) or virtual lock (Windows)



Lesson 4

# Review - Server Configuration Best Practices



# Summary

- For best performance, choose **SSD over spinning disks**
- Local disks are preferred - the software provides HA
- In general, choose **large** machines over **x-large** machines
- Elasticsearch has very good JVM defaults, so avoid changing them or be careful
- When configuring the JVM heap, leave as much memory to the filesystem cache as possible

# Quiz

- 1. True or False.** SAN storage is preferred over local disks to provide high availability of data.
- 2. True or False.** It is a good idea to separate the **transport** and **HTTP** traffic over different network interfaces.
- 3. True or False.** It is a good idea to change the JVM defaults to accommodate your needs.

